

Customer Segmentation, Lifetime Value Prediction, and Churn Analysis of a Turkish E-Commerce Platform

Ivan Howes & Maria Smelkov

1. Abstract

This report analyzes a simulated Turkish e-commerce dataset containing 17,049 transactions across 5,000 unique customers spanning January 2023 to March 2024. During initial exploration, we found evidence that the dataset was synthetically generated discount rates, delivery times, monthly order volumes, and satisfaction ratings all exhibit unusually low variation. As a result, we aggregate data to the customer level, where more meaningful behavioral variation emerges.

We address three questions: (1) Can customers be grouped into meaningful types based on purchasing behavior? (2) What characteristics are most associated with higher total spending? (3) Can we identify customers likely to stop purchasing? To answer these, we apply three methods. First, RFM (Recency, Frequency, Monetary) analysis combined with K-Means clustering segments customers into four groups: Champions, Loyal, At Risk, and Inactive. Second, an Ordinary Least Squares regression model predicts Customer Lifetime Value. Third, a logistic regression model to predict customer churn defined as no purchase in 180 or more days.

Champions (roughly 24% of customers) have an average predicted lifetime spend of £9,742, while Inactive customers average just £787. Churn rates range from 3.6% among Champions to 51.8% among Inactive customers. These findings offer practical guidance for targeted marketing, retention strategies, and customer resource allocation.

2. Introduction

Online retail generates large volumes of customer transaction data, creating opportunities to understand purchasing behavior in detail. Businesses that can identify high-value customers, predict future spending, and detect early signs of dropout are better positioned to allocate marketing resources efficiently and reduce revenue loss from churn.

The RFM framework is an established approach to customer segmentation. It captures the three behavioral dimensions most predictive of future purchasing. Pairing RFM with K-Means clustering allows distinct customer types to emerge from the data rather than being imposed by arbitrary cutoffs.

A key challenge in this project is that the dataset is synthetically generated. Discount rates are nearly uniform across product categories, delivery times barely vary across cities, monthly order volumes stay nearly constant throughout the year, and satisfaction ratings cluster tightly regardless of context. This means transaction-level relationships are unlikely to reflect meaningful real-world behavior. Our response is to shift analysis to the customer level, aggregating each customer's history into summary statistics where genuine variation does exist.

Research Questions

1. Can customers be grouped into meaningful types based on purchasing behavior?
2. What customer characteristics are most associated with higher total spending?
3. Can we identify customers likely to stop purchasing?

Key Results

- K-Means with $K=4$ produced four segments: Champions (1,188), Loyal (757), At Risk (1,864), and Inactive (1,191).
- OLS regression predicted log-transformed lifetime spend with $R^2 = 0.585$ and $MAE = 0.676$.
- A balanced logistic regression achieved 51.6% churn sensitivity vs. near-zero for the unweighted model.
- Overall churn rate is 26.2%, ranging from 0.0% (Loyal) to 51.8% (Inactive).

Document Structure

Section 2 covers methodology. Section 3 covers data sourcing, cleaning, and exploration. Section 4 presents model results. Section 5 concludes with findings and future directions. Section 6 contains references.

3. Methodology

3.1 RFM Analysis

RFM summarizes each customer's purchase history along three dimensions:

- **Recency (R):** Days since the customer's most recent purchase (relative to the snapshot date — the latest date in the dataset). Lower = more recent.
- **Frequency (F):** Total number of distinct orders placed.
- **Monetary (M):** Total amount spent across all orders.

3.2 K-Means Clustering

K-Means is an unsupervised machine learning algorithm that partitions observations into K clusters by minimizing within-cluster variation. We describe it here in detail.

The algorithm minimizes the within-cluster sum of squared distances:

$$\arg \min_C \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

where C_k is the set of points in cluster k and μ_k is the centroid (mean) of that cluster.

Algorithm steps:

1. Pick K random starting points called centroids, one per cluster.
2. Assign every customer to their nearest centroid.
3. Move each centroid to the average position of its assigned customers.
4. Repeat steps 2–3 until the assignments stop changing.

Because the starting points are random, the algorithm can sometimes land on a poor solution, so we run it 10 times with different starting points and keep the best result.

Pre-processing: K-Means groups points by distance, so features on a larger scale will dominate. For example, Monetary values in the thousands would completely overshadow Frequency values of 1–10. To fix this, we first apply a log transformation to reduce the influence of extreme values, then rescale all three features to the same range.

Choosing K: We don't know in advance how many customer groups exist. We try $K = 2$ through 10 and measure how tight the clusters are at each value. Plotting this gives us an "elbow" shape, after a certain K, adding more clusters barely improves the tightness. That elbow point is our optimal K.

3.3 OLS Regression (CLV Prediction)

We predict Customer Lifetime Value using Ordinary Least Squares regression. The target is log-transformed total spend, $\log(1 + \text{Monetary})$, to reduce skewness. Since Recency can have values of 0, we are using $\log(1 + \text{Monetary})$ to shift everything up by 1 therefore the smallest possible input is 1.

$$\log(1 + \text{Monetary}_i) = \beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip} + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$

Categorical variables are coded as either a 1 (True) or 0 (False) to avoid multicollinearity. The model is trained on 80% of customers and tested on the remaining 20% using two metrics: R^2 (how much spending variation the model explains) and MAE (average prediction error).

3.4 Logistic Regression (Churn Prediction)

Churn is defined as Recency ≥ 180 days. We model churn probability as:

$$P(\text{Churned}_i = 1) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip})}}$$

Data leakage: Recency and Frequency are excluded from the feature set because, of course, Recency is directly correlated with Frequency.

Class imbalance: The churn rate is 26.2%, so a naive model can reach 73.8% accuracy by simply predicting "not churned" for everyone. We compare:

- **Standard:** No adjustment. Nearly all predictions are "not churned."
- **Balanced** (`class_weight='balanced'`): Upweights churned observations by applying a higher penalty when Churned customers are not correctly identified. This way, the model pays more attention to the minority churned class.

We report accuracy, sensitivity, and specificity.

4. Data

4.1 Source

The dataset is the *E-Commerce Customer Behavior and Sales Analysis* dataset (Tuygur, 2024), downloaded from Kaggle via the `kagglehub` API. It simulates Turkish e-commerce customer behavior and contains both purchase and web session data across 17,049 transactions from 5,000 customers (January 2023 – March 2024).

4.2 Variables

Variable	Type	Description
Order_ID	String	Unique order identifier
Customer_ID	String	Unique customer identifier
Date	String → datetime	Order date
Age	Integer	Customer age
Gender	Categorical	Customer gender
City	Categorical	Customer city
Product_Category	Categorical	Product category
Unit_Price	Float	Price per unit (£)
Quantity	Integer	Units purchased
Discount_Amount	Float	Discount applied (£)
Total_Amount	Float	Final transaction value (£)
Payment_Method	Categorical	Payment type
Device_Type	Categorical	Mobile, Desktop, or Tablet
Session_Duration_Minutes	Integer	Browsing session length (minutes)

Variable	Type	Description
Pages_Viewed	Integer	Pages visited per session
Is_Returning_Customer	Boolean	Whether customer has shopped before
Delivery_Time_Days	Integer	Days from order to delivery
Customer_Rating	Integer	Post-purchase rating (1–5)

4.3 Synthetic Data

Several variables show signs of artificial generation. This means individual transaction-level relationships are not reliable. Our solution is to work at the customer level, aggregating each customer's order history into summary statistics where meaningful variation does exist. Conclusions are **descriptive and exploratory**, not generalizable to real-world Turkish e-commerce.

4.4 Data Preparation

Shape: (17049, 18)

Data types:

```

Order_ID          str
Customer_ID       str
Date              str
Age               int64
Gender            str
City              str
Product_Category  str
Unit_Price        float64
Quantity          int64
Discount_Amount   float64
Total_Amount      float64
Payment_Method     str
Device_Type       str
Session_Duration_Minutes int64
Pages_Viewed      int64
Is_Returning_Customer bool
Delivery_Time_Days int64
Customer_Rating   int64
dtype: object

```

	Order_ID	Customer_ID	Date	Age	Gender	City	Product_Category	Unit_Price	G
0	ORD_000001-1	CUST_00001	2023-05-29	40	Male	Ankara	Books	29.18	
1	ORD_000001-2	CUST_00001	2023-10-12	40	Male	Ankara	Home & Garden	644.40	
2	ORD_000001-3	CUST_00001	2023-12-05	40	Male	Ankara	Sports	332.82	
3	ORD_000002-1	CUST_00002	2023-05-11	33	Male	Istanbul	Food	69.30	
4	ORD_000002-2	CUST_00002	2023-06-16	33	Male	Istanbul	Beauty	178.15	
5	ORD_000003-1	CUST_00003	2023-02-27	42	Male	Konya	Toys	198.28	
6	ORD_000003-2	CUST_00003	2024-01-03	42	Male	Konya	Home & Garden	526.85	
7	ORD_000004-1	CUST_00004	2024-02-13	53	Male	Izmir	Fashion	96.20	
8	ORD_000005-1	CUST_00005	2023-03-16	32	Male	Ankara	Home & Garden	533.67	
9	ORD_000005-2	CUST_00005	2023-06-12	32	Male	Ankara	Toys	73.06	

Missingness and Duplicates

Missing values per column:

Order_ID	0
Customer_ID	0
Date	0
Age	0
Gender	0
City	0
Product_Category	0
Unit_Price	0
Quantity	0
Discount_Amount	0
Total_Amount	0
Payment_Method	0
Device_Type	0
Session_Duration_Minutes	0
Pages_Viewed	0
Is_Returning_Customer	0
Delivery_Time_Days	0
Customer_Rating	0

dtype: int64

Duplicate rows: 0

No missing values or duplicate rows were found.

Type Correction

```
In [41]: # Convert Date from string to datetime
df['Date'] = pd.to_datetime(df['Date'])
print("Date range:", df['Date'].min(), "to", df['Date'].max())
```

Date range: 2023-01-01 00:00:00 to 2024-03-25 00:00:00

Restructuring: Transaction → Customer Level

Customers: 5000

	Customer_ID	Age	Avg_Session_Duration	Avg_Pages_Viewed	Avg_Customer_Rating	Avg_I
0	CUST_00001	40	14.333333	9.000000	3.333333	
1	CUST_00002	33	15.000000	10.000000	4.000000	
2	CUST_00003	42	10.500000	8.500000	3.500000	
3	CUST_00004	53	16.000000	15.000000	5.000000	
4	CUST_00005	32	12.666667	9.333333	3.666667	

RFM Computation


RFM shape: (5000, 4)

	Customer_ID	Recency	Frequency	Monetary
0	CUST_00001	111	3	2199.63
1	CUST_00002	283	2	809.90
2	CUST_00003	82	2	3030.81
3	CUST_00004	41	1	383.22
4	CUST_00005	278	3	2422.73
5	CUST_00006	211	3	1245.83
6	CUST_00007	80	1	1331.55
7	CUST_00008	168	3	5602.29
8	CUST_00009	38	1	470.86
9	CUST_00010	21	3	1421.16


	Recency	Frequency	Monetary
count	5000.00	5000.00	5000.00
mean	124.72	3.41	4355.81
std	108.05	2.18	5274.55
min	0.00	1.00	14.28
25%	38.00	2.00	889.41
50%	93.00	3.00	2493.93
75%	184.00	5.00	5796.30
max	449.00	10.00	50628.15

4.5 Exploratory Data Analysis

RFM Distributions



The raw RFM distributions are all right-skewed; a small number of customers have very high spend or frequency. After log transformation and standardization, distributions become approximately symmetric.




After log-transforming the data to reduce the skew, our data is prepared for the K-means model.

5. Results

5.1 Customer Segmentation (RFM + K-Means)

Choosing K — Elbow Method



Inertia drops sharply from K=2 to K=4, then flattens supporting K=4 as the optimal choice.

Fitting the Final Model (K=4)

Customers per cluster:

Cluster

0 757

1 1191

2 1864

3 1188

Name: count, dtype: int64

Cluster profiles:

	Recency	Frequency	Monetary
Cluster			
0	12.49	4.08	4866.98
1	202.39	1.44	473.17
2	155.73	2.63	3633.03
3	69.69	6.18	9056.59

K-Means partitions all 5,000 customers into four groups by placing each customer into the cluster whose center (centroid) is closest to their RFM values. The table above shows the average Recency, Frequency, and Monetary value at the center of each cluster. Customers whose purchasing behavior most resembles these averages were assigned to that group.

Below we assign titles to these clusters that represent the type of customers within them and update the dataframe with a new segment column:


Segment
 At Risk 1864
 Inactive 1191
 Champions 1188
 Loyal 757
 Name: count, dtype: int64

	Customer_ID	Recency	Frequency	Monetary	Cluster	Segment
0	CUST_00001	111	3	2199.63	2	At Risk
1	CUST_00002	283	2	809.90	1	Inactive
2	CUST_00003	82	2	3030.81	2	At Risk
3	CUST_00004	41	1	383.22	1	Inactive
4	CUST_00005	278	3	2422.73	2	At Risk
5	CUST_00006	211	3	1245.83	2	At Risk
6	CUST_00007	80	1	1331.55	1	Inactive
7	CUST_00008	168	3	5602.29	2	At Risk
8	CUST_00009	38	1	470.86	1	Inactive
9	CUST_00010	21	3	1421.16	0	Loyal

Segment interpretations:

Segment	Avg. Recency	Avg. Frequency	Avg. Monetary	Description
Champions	69.7 days	6.2 orders	£9,057	High spend, high frequency, recent. Best customers.
Loyal	12.5 days	4.1 orders	£4,867	Very recently active, buy regularly. Prime to become Champions.
At Risk	155.7 days	2.6 orders	£3,633	Decent spend history but disengaging. Target for re-engagement.
Inactive	202.4 days	1.4 orders	£473	Low spend, low frequency, haven't bought in 6+ months.

This answers Research Question 1: **yes, customers group into four distinct and interpretable behavioral types.**



We see that the biggest segment of customers are "At Risk" (~1750), the smallest segment is "Loyal" (~750). This business should aim marketing strategies to target these "At Risk" customers and reduce churn.

5.2 Customer Lifetime Value Prediction

(5000, 17)


	Customer_ID	Recency	Frequency	Monetary	Cluster	Segment	Age	Avg_Session_Duratio
0	CUST_00001	111	3	2199.63	2	At Risk	40	14.33333
1	CUST_00002	283	2	809.90	1	Inactive	33	15.00000
2	CUST_00003	82	2	3030.81	2	At Risk	42	10.50000
3	CUST_00004	41	1	383.22	1	Inactive	53	16.00000
4	CUST_00005	278	3	2422.73	2	At Risk	32	12.66666

We combine our three customer-level tables, rfm, numeric_agg, and cat_agg, into a single feature table to use in the CLV regression model.

Fitting the Regression Model for Customer Lifetime Value

R² score: 0.585
MAE: 0.676

The model explains ~59% of variation in log customer spend. This is a reasonable result for a linear model on behavioral data.




The strongest positive predictors of lifetime spend are product categories like Electronics, Home & Garden, and Sports along with purchase Frequency. Food is the strongest negative predictor, followed by Tablet device usage. Most payment method and device type effects are small and only marginally significant.

The dominance of Electronics, Home & Garden, and Sports makes intuitive sense, as these categories typically involve high-ticket purchases that drive total spend up significantly. Conversely, Food customers likely make frequent but low-value purchases, pulling their lifetime monetary value down relative to other categories.

Predicted CLV by Segment

To check the validity of our K-means and Regression models, we look at the lifetime spending by customer segment.

```
Segment
Champions    9742.30
Loyal        4758.07
At Risk      2516.04
Inactive      787.00
Name: Predicted_CLV, dtype: float64
```



The predicted CLV values align closely with the RFM cluster profiles, validating both approaches. This confirms our answer to Research Question 2: **expensive product categories and order frequency are the strongest drivers of lifetime spend**, with Champions averaging £9,742 vs. £787 for Inactive customers.

5.3 Churn Prediction

Fitting the Regression Model for Churn Prediction

Churn counts:

Churned

0 3692

1 1308

Name: count, dtype: int64

Overall churn rate: 26.2%

Out[54]:

	Customer_ID	Recency	Frequency	Monetary	Cluster	Segment	Predicted_CLV	Churned
0	CUST_00001	111	3	2199.63	2	At Risk	1184.354695	0
1	CUST_00002	283	2	809.90	1	Inactive	895.052796	1
2	CUST_00003	82	2	3030.81	2	At Risk	1907.795274	0
3	CUST_00004	41	1	383.22	1	Inactive	1076.345236	0
4	CUST_00005	278	3	2422.73	2	At Risk	1330.591549	1

We add a new boolean variable to our dataframe, "Churned", that defines customers that have not ordered in 180 days or more as a churned customer.

Logistic Regression — Standard vs. Balanced


Accuracy: 0.757
Sensitivity: 0.004
Specificity: 1.000

The standard model achieves 75.7% accuracy, but this is misleading because it reaches this by predicting 'not churned' for nearly every customer, resulting in a sensitivity of just 0.4%. For churn detection, this is essentially useless since correctly identifying at-risk customers is the goal.

Accuracy: 0.534
Sensitivity: 0.516
Specificity: 0.540

To address this, we refit the model with `class_weight='balanced'`, which penalizes misclassifying churned customers more heavily. This drops accuracy to 53.4%, but raises sensitivity to 51.6% which is a far more useful result for identifying customers at risk of churning. Missing a customer who is about to churn is more costly than accidentally targeting one who wasn't going to, hence why sensitivity matters more than overall accuracy here.

Logistic Regression Coefficients (Balanced Model)



Lower session engagement (fewer pages viewed, shorter session duration) is associated with higher churn risk. Notably, customers who pay by credit card show a strong negative association with churn, suggesting that the commitment of linking a card correlates with longer-term purchasing behavior. Similarly, mobile device users are less likely to churn, possibly reflecting habitual app-based shopping. This answers Research Question 3: **yes, customers likely to stop purchasing can be identified.**

6. Conclusion

Findings

This project applied RFM clustering, OLS regression, and logistic regression to customer-level data from a simulated Turkish e-commerce platform. Despite limitations from synthetic data, the customer-level aggregation revealed genuine behavioral variation that supported meaningful modeling.

By research question:

1. **Can customers be grouped into meaningful types?** Yes. K-Means with K=4 produced four well-separated, interpretable segments, Champions, Loyal, At Risk, and Inactive, that differ clearly in recency, frequency, spend, and churn rates.
2. **What drives higher lifetime spending?** Purchase frequency and product category are the strongest predictors. Electronics, Home & Garden, and Sports customers spend significantly more due to the high-price nature of those purchases, while Food customers show the opposite effect. The CLV regression ($R^2 = 0.585$) confirms these patterns, and predicted spend aligns closely with the unsupervised cluster structure, with Champions averaging £9,742 vs. £787 for Inactive customers.
3. **Can we identify customers likely to churn?** The standard logistic regression is essentially useless for this task, but the balanced model improves sensitivity to 51.6%, identifying roughly half of at-risk customers, a meaningful result for targeted advertisement. Credit card users and mobile shoppers show lower churn risk, while lower session engagement is the strongest behavioral signal of impending churn.

Business Implications

- **Champions** should be retained with loyalty rewards or early product access.
- **Loyal** customers are the best candidates to move into the Champions segment with targeted incentives.
- **At Risk** customers warrant a re-engagement campaign before they become Inactive.
- **Inactive** customers have low conversion potential and should receive minimal marketing spend.

From the CLV model, the business should prioritize acquiring and retaining Electronics, Home & Garden, and Sports customers, as these categories are the strongest drivers of lifetime spend.

From the churn model, the business should pay particular attention to customers showing declining session engagement (fewer pages viewed and shorter browsing sessions) is a warning signs of churn. Encouraging credit card linking and mobile app adoption may also help reduce churn, as both are associated with lower dropout rates.

Limitations

The synthetic nature of the data is the primary limitation i.e. patterns may not reflect real consumer behavior. Additionally, the 180-day churn threshold is an assumption; different cutoffs would produce different results.

Future Work

Based on our results some natural follow-up questions emerge:

- What distinguishes the At Risk customers who eventually recover vs. those who become Inactive? Are there early warning signs that separate them?
- The CLV model explained ~59% of spending variation, meaning 41% is unexplained. What other factors might account for this? Seasonal purchasing patterns, promotional responses, or life events?
- Given that product category was such a dominant CLV predictor, it would be interesting to ask whether a customer's first purchase category is predictive of their long-term spending trajectory.
- We defined churn as 180+ days of inactivity, but this is arbitrary. How sensitive are churn model results to different cutoff choices (say 90 days or 365 days)?
- The model only identifies customers who have already effectively churned. Could we build an earlier warning system that flags customers showing declining engagement before they hit the 180-day threshold?

7. References

Tuygur, U. (2024). E-commerce customer behavior and sales analysis [Dataset]. Kaggle. <https://www.kaggle.com/datasets/umuttuygurr/e-commerce-customer-behavior-and-sales-analysis-tr>

IBM. (n.d.). What is k-means clustering? IBM Think. <https://www.ibm.com/think/topics/k-means-clustering>

Gusarova, M. (2023, February 5). How to improve logistic regression in imbalanced data with class weights. Medium. <https://medium.com/@data.science.enthusiast/how-to-improve-logistic-regression-in-imbalanced-data-with-class-weights-1693719136aa>

